

JULIAN ESTEBAN RINCÓN RODRÍGUEZ

Computer Science & Artificial Intelligence

Bogotá, Colombia • julianer2002@gmail.com • +57 311 828 3342 • github.com/Julian-Rincon

PROFESSIONAL PROFILE

8th-semester Computer Science & AI student at Universidad Sergio Arboleda, specialized in data science, machine learning, and intelligent systems. Hands-on experience building data pipelines, training predictive models (XGBoost, LightGBM, neural networks), cloud infrastructure on AWS, and local AI applications. Proven ability to take projects beyond the academic setting and deliver production-grade, documented work.

EDUCATION

Universidad Sergio Arboleda — Bogotá, Colombia

2021 – Present

B.Sc. Computer Science & Artificial Intelligence | 8th semester of 9

Current courses (8th sem.): Big Data & Data Engineering, Machine Learning, Simulation & Modeling, Neural Networks — Classification (Elective), Engineering Communication, Ethics & Social Responsibility.

FEATURED PROJECTS

Jarvis IronMan — Dual-Brain Local AI Assistant

2025 – Present

Python 3.11 · Ollama (fast local routing) · Gemini API (complex reasoning) · Whisper STT · Piper TTS · Playwright · ZeroMQ · ChromaDB

- Engineered dual-brain inference architecture: Ollama handles sub-200ms local intent routing; Gemini API resolves complex multi-step reasoning — zero cloud dependency for core pipeline.
- Architected multiprocessing IPC layer via ZeroMQ isolating STT, LLM, and TTS processes; 12 active skills including Gmail/Calendar OAuth2, Playwright browser automation, and Telegram remote control.
- Deployed wake-word detection (openwakeword) with STT fallback; custom tkinter HUD rendering system state, active skill, and latency metrics in real time.
- Engineering ChromaDB vector memory for cross-session context persistence; autonomous meeting agent in development (system audio STT → Gemini reasoning → Piper TTS response).

SAVI v2 — Autonomous Real Estate Valuation System

2026

Python · XGBoost · K-Means · MDP · Q-Learning · Deep Q-Network · PyTorch · scikit-learn

- Architected full RL decision pipeline: MDP Value Iteration (convergence 259 iter, $\gamma=0.95$) → Q-Learning tabular (8K episodes, ϵ -greedy decay) → Double DQN (PyTorch, 150 epochs, replay buffer 20K, Adam $\text{lr}=3\text{e-}4$).
- Consensus policy across all three agents autonomously APPROVE/REVIEW/REJECTs valuations based on expected financial risk; XGBoost + cluster feature achieves $R^2=0.9609$ on 20,203 properties (2006–2024).
- Delivered IEEE-format technical paper, interactive HTML demo on GitHub Pages, and full reproducible Python pipeline — published to production portfolio.

Project Dogma — Stochastic Social Propagation Engine

2026

TypeScript · React · Dynamic Markov Chains · Hidden Markov Model · Nash Equilibrium · Monte Carlo

- Architected a mathematical simulation engine for ideological propagation across urban social graphs: Dynamic Markov Chains model agent state transitions (Neutral → Believer → Fanatic) with per-agent solitude, skepticism, and charisma variables modifying transition matrices each tick.
- Modeled authority behavior via Hidden Markov Model (Forward Algorithm): four latent states inferred from a single observable signal (Heat entropy), generating probabilistic threat cycles without exposing internal state.
- Resolved doctrinal conflicts via 200 Nash Equilibrium mixed strategy (55%/45%), eliminating dominant pure strategies; resource economy governed by Verhulst logistic differential equation (carrying capacity K). Validated across 500 Monte Carlo simulations: $\mu=152$ ticks, conversion 5.42 agents/tick ($\sigma=0.39$), retention ratio 38:1.

ShopStream — End-to-End AWS Big Data Pipeline

2026

[Python](#) · [PySpark](#) · [AWS Lambda](#) · [AWS EMR](#) · [AWS Glue](#) · [RDS PostgreSQL](#) · [Flask](#) · [Zappa](#) · [GitHub Actions](#)

- Engineered production-grade data pipeline processing 2.5M synthetic behavioral events (page_view, click, search, product_view, cart_event) partitioned by date across S3 raw and processed zones.
- Deployed S3-triggered Lambda validator with CloudWatch custom metrics and automatic quarantine of invalid records; EMR/PySpark job computes 6 business KPIs including bounce rate, conversion funnel, and anomaly detection.
- Delivered Flask REST API on AWS Lambda via Zappa exposing DWH analytics endpoints; 86% test coverage with pytest; CI/CD via GitHub Actions (lint, test, deploy, smoke-test); RDS security group hardened from 0.0.0.0/0 to /32 CIDR.

Project Dogma — Stochastic Social Propagation Engine

2026

[TypeScript](#) · [React](#) · [Dynamic Markov Chains](#) · [HMM](#) · [Verhulst ODE](#) · [Nash Equilibrium](#) · [Monte Carlo](#)

- Architected stochastic social propagation simulator modeling ideological spread across urban agent graphs; Dynamic Markov Chains govern state transitions (Neutral → Believer → Fanatic) modulated by per-agent solitude, skepticism, and charisma variables.
- Implemented HMM (Forward algorithm) for authority behavior inference from latent states via observable Heat entropy signal; Verhulst logistic differential equation governs resource carrying capacity K, generating S-curve propagation dynamics.
- Resolved doctrinal conflict via 2×2 Nash Equilibrium mixed strategy (55%/45%); validated across 500 Monte Carlo simulations: $\mu=152$ ticks to convergence, conversion 5.42 agents/tick ($\sigma=0.39$), retention ratio 38:1.

Big Data Platform — AWS (Chinook / Sakila / Credit)

2026

[AWS EC2](#) · [RDS \(PostgreSQL / MySQL\)](#) · [Athena](#) · [S3](#) · [Glue](#) · [Terraform](#) · [React](#) · [FastAPI](#)

- Fullstack on Chinook DB: React/Vite frontend + FastAPI backend on separate EC2 instances, PostgreSQL on RDS; 17-section technical document.
- Hive-partitioned external table (credit2) on Athena + S3; resolved case-sensitivity issue in partitions with manual ALTER TABLE.
- NYC Yellow Taxi dataset (94M rows, 12 months) on AWS EMR with PySpark; resolved schema inconsistency across Parquet files using explicit StructType and per-file normalization with unionByName.

Neural Networks & RL Labs

2026

[PyTorch](#) · [TensorFlow](#) · [CUDA](#) · [CIFAR-10](#) · [OpenAI Gym](#)

- Reinforcement learning: Q-learning (FrozenLake), REINFORCE (CartPole), and DQN with replay buffer.
- CNN labs: manual convolution, SmallCNN with dropout and weight decay trained on CIFAR-10; professional Word reports.

Ames Housing ML Pipeline

2026

[Python](#) · [XGBoost](#) · [LightGBM](#) · [scikit-learn](#) · [Pandas](#) · [Jupyter Notebook](#)

- Best model: LightGBM $R^2 = 0.75$, MAE $\approx \$23,600$ on $\sim 20,000$ records (2006–2024 dataset).
- Complete pipeline: EDA, feature engineering, K-Means and hierarchical clustering, comparison of 6 algorithms.

Custom Programming Language

2023

[Python](#) · [ANTLR4](#) · [Tokenization](#) · [Interpretation](#)

- Designed and implemented a Python-inspired language with custom tokenizer, lexer, parser, and visitor-pattern execution engine using ANTLR4.

TECHNICAL SKILLS

Languages

Python (intermediate–advanced), Bash Scripting, JavaScript (basic), SQL

ML / AI

scikit-learn, XGBoost, LightGBM, PyTorch, TensorFlow, Whisper, Ollama

Big Data & Cloud

AWS (EC2, RDS, S3, Athena, Glue, EMR), Terraform, Git / GitHub, Docker (basic)

Data & ETL

Pandas, NumPy, Matplotlib, openpyxl, Power BI, Jupyter Notebook, PySpark

Automation

Playwright, pywinauto, Telegram Bot API, Gmail / Calendar OAuth2

Other

ANTLR4, CUDA, NetworkX, ChromaDB, ZeroMQ, tkinter, pytest

ACHIEVEMENTS & RECOGNITION

1st Place — Wisibilízalas International Competition

2019

Universidad Pompeu Fabra · Barcelona, Spain · Shakespeare School, Quito, Ecuador · Senior Category

- 1st place in international competition organized by Universidad Pompeu Fabra, representing Shakespeare School (Ecuador) in the Senior category at age 17.
- Winning project: website highlighting the careers of women in technology and ICT, competing against teams from Latin America and Europe.
- Project available at: sites.google.com/view/warmitics

Independent Project — Jarvis IronMan

Self-directed development outside the academic curriculum

- Fully functional local AI voice assistant built from scratch, without relying on cloud APIs. Personal initiative beyond the academic curriculum.

CERTIFICATIONS

Database Design — Oracle Academy — August 2024

Database Programming with SQL — Oracle Academy — December 2024

AWS Academy Graduate — Data Engineering — Amazon Web Services Academy — May 2026 — 40 hours · Credly: credly.com/go/UGp1CUEB

LANGUAGES

Spanish: Native • English: B2+ (technical and academic use — documentation, presentations, reports)

Available for internships, freelance projects, or collaborations in AI, Data Science, and Software Development.